

# Heart Disease Risk Prediction

## Objective

The goal of this project is to build a machine learning model that predicts the likelihood of heart disease based on patient health metrics.

### Importing required libraries for the Machine Learning Workflow

```
In [2]: import pandas as pd  
import numpy as np  
import seaborn as sns  
import matplotlib.pyplot as plt
```

### Loading the Dataset

```
In [3]: df = pd.read_csv('heart_data.csv')  
df.head(5)
```

```
Out[3]:
```

	age	sex	chest pain type	resting bp s	cholesterol	fasting blood sugar	resting ecg	max heart rate	exercise angina	oldpeak s
0	40	1	2	140	289	0	0	172	0	0.0
1	49	0	3	160	180	0	0	156	0	1.0
2	37	1	2	130	283	0	1	98	0	0.0
3	48	0	4	138	214	0	0	108	1	1.5
4	54	1	3	150	195	0	0	122	0	0.0

```
In [4]: df.tail(5)
```

Out[4]:

	age	sex	chest pain type	resting bp s	cholesterol	fasting blood sugar	resting ecg	max heart rate	exercise angina	oldpeak
1185	45	1	1	110	264	0	0	132	0	1.2
1186	68	1	4	144	193	1	0	141	0	3.4
1187	57	1	4	130	131	0	0	115	1	1.2
1188	57	0	2	130	236	0	2	174	0	0.0
1189	38	1	3	138	175	0	0	173	0	0.0

## Checking numbers of rows and columns in data set

In [5]: df.shape

Out[5]: (1190, 12)

In [6]: df.describe()

	age	sex	chest pain type	resting bp s	cholesterol	fasting blood sugar
count	1190.000000	1190.000000	1190.000000	1190.000000	1190.000000	1190.000000
mean	53.720168	0.763866	3.232773	132.153782	210.363866	0.213445
std	9.358203	0.424884	0.935480	18.368823	101.420489	0.409912
min	28.000000	0.000000	1.000000	0.000000	0.000000	0.000000
25%	47.000000	1.000000	3.000000	120.000000	188.000000	0.000000
50%	54.000000	1.000000	4.000000	130.000000	229.000000	0.000000
75%	60.000000	1.000000	4.000000	140.000000	269.750000	0.000000
max	77.000000	1.000000	4.000000	200.000000	603.000000	1.000000

## Checking for Missing or Zero Values

In [7]: (df == 0.0).sum()

```
Out[7]: age          0
         sex         281
         chest pain type    0
         resting bp s      1
         cholesterol     172
         fasting blood sugar 936
         resting ecg       684
         max heart rate    0
         exercise angina   729
         oldpeak           455
         ST slope          1
         target            561
         dtype: int64
```

## Exploring Data Types and Target Distribution

```
In [8]: print(df.dtypes)
```

```
age          int64
sex          int64
chest pain type  int64
resting bp s    int64
cholesterol    int64
fasting blood sugar int64
resting ecg     int64
max heart rate  int64
exercise angina int64
oldpeak        float64
ST slope        int64
target          int64
dtype: object
```

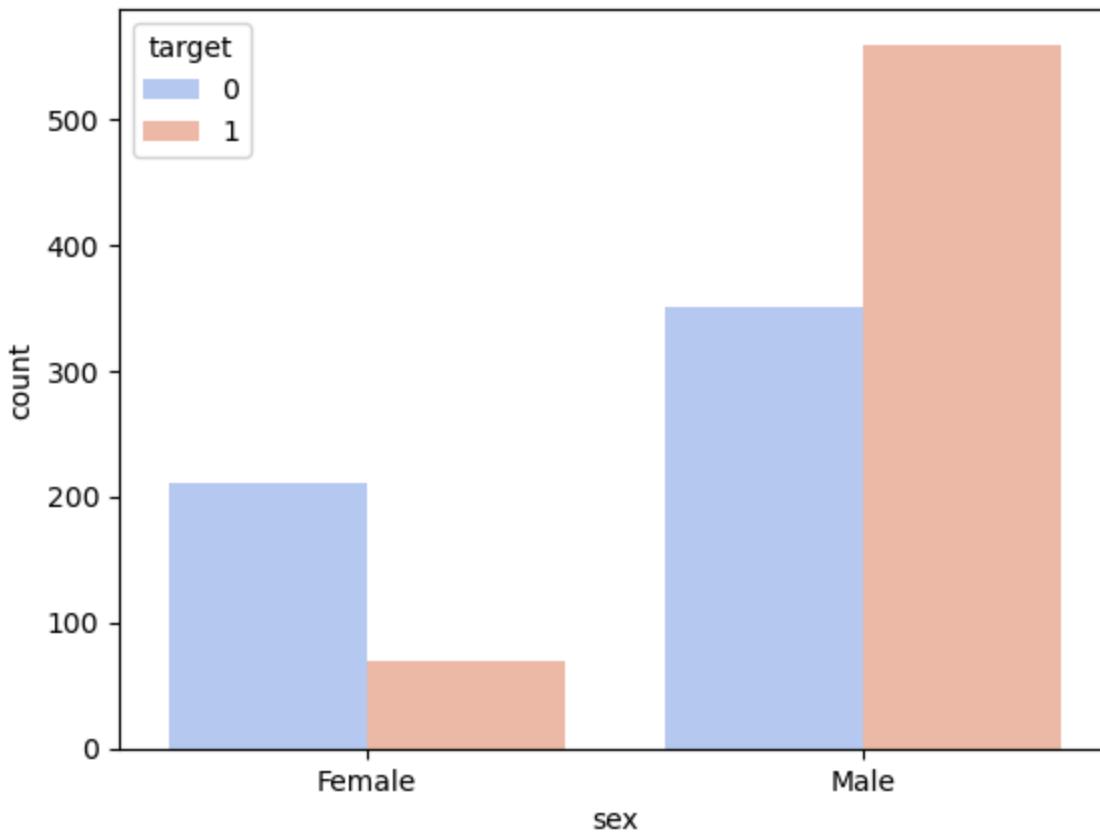
```
In [9]: print(df['target'].value_counts())
```

```
target
1    629
0    561
Name: count, dtype: int64
```

Visualizes how categorical features relate to heart disease. For example, the sex countplot shows the number of males and females with or without heart disease, helping identify patterns and important features for prediction.

```
In [10]: sns.countplot(x='sex', hue='target', data=df, palette='coolwarm')
plt.title("Sex vs Heart Disease")
plt.xticks([0,1], ['Female','Male'])
plt.show()
```

## Sex vs Heart Disease

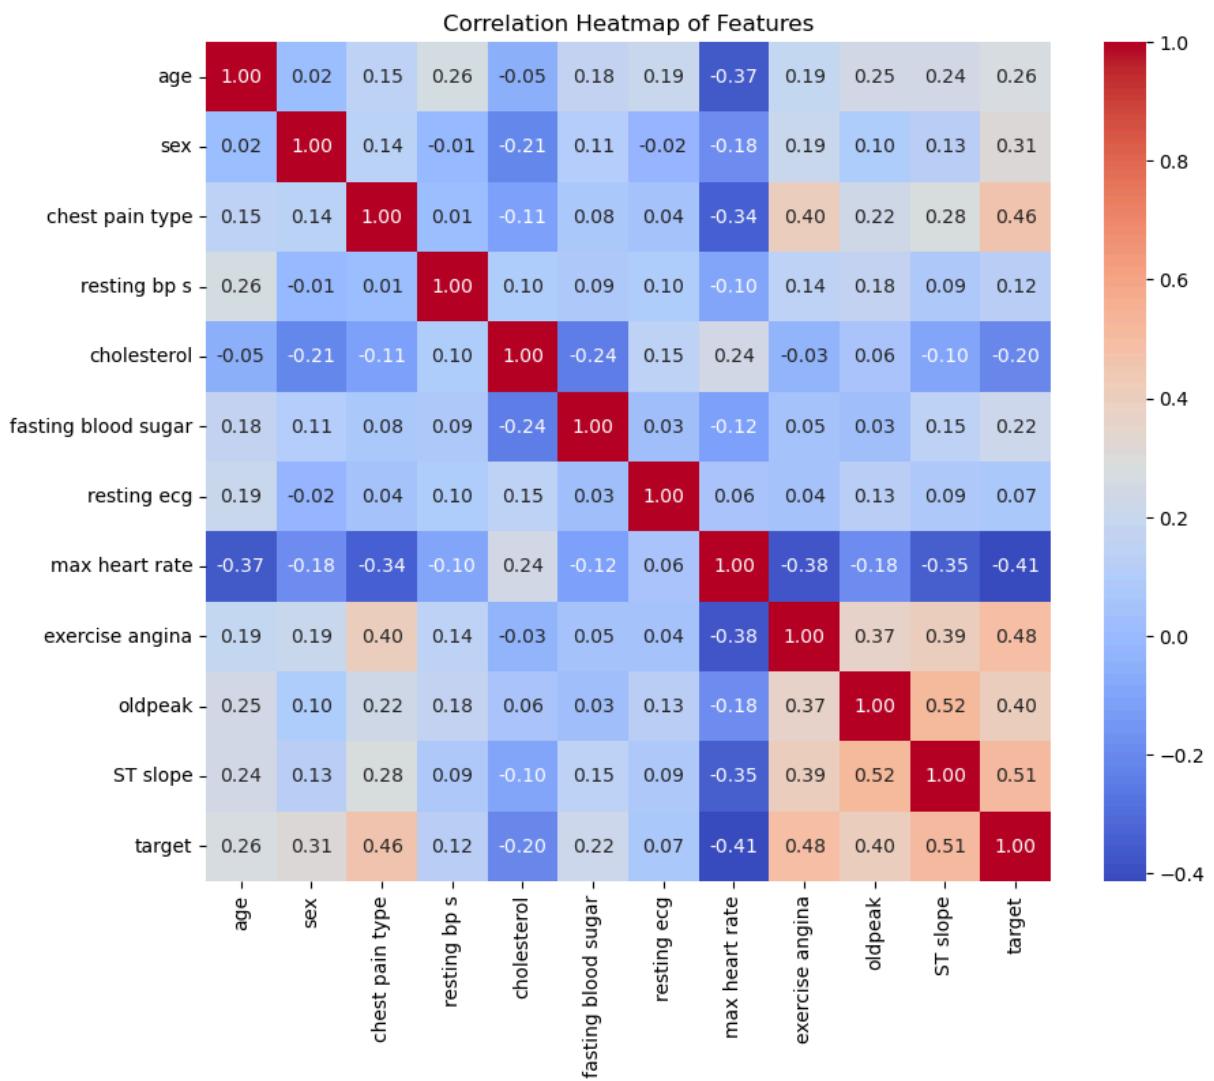


## Correlation Heatmap

Shows how features are correlated with each other and with the target.

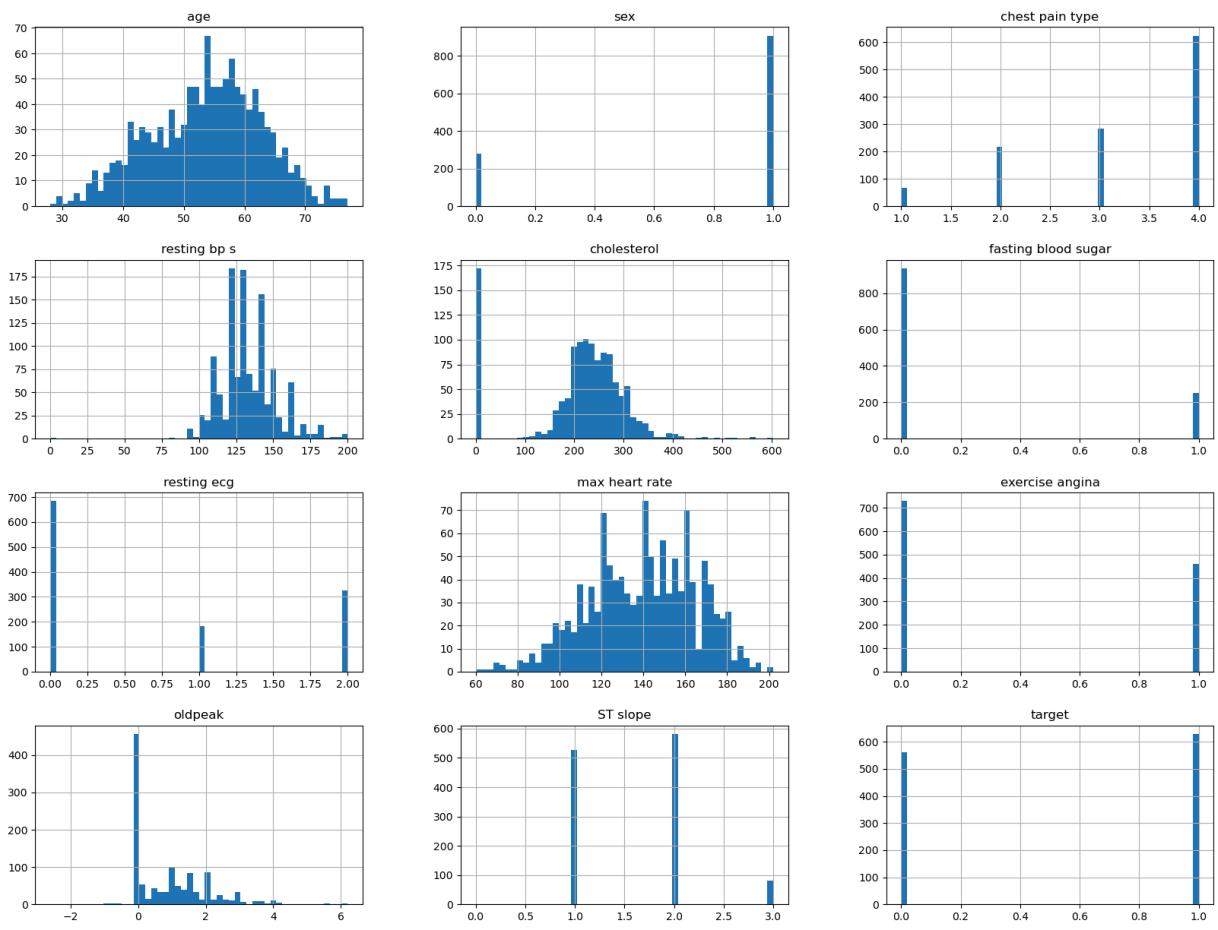
Helps in understanding which features may be more important for prediction.

```
In [11]: plt.figure(figsize=(10,8))
sns.heatmap(df.corr(), annot=True, cmap="coolwarm", fmt=".2f")
plt.title("Correlation Heatmap of Features")
plt.show()
```



```
In [12]: df.hist(bins = 50, figsize = (20,15))
```

```
Out[12]: array([[<Axes: title={'center': 'age'}>, <Axes: title={'center': 'sex'}>,
   <Axes: title={'center': 'chest pain type'}>],
  [<Axes: title={'center': 'resting bp s'}>,
   <Axes: title={'center': 'cholesterol'}>,
   <Axes: title={'center': 'fasting blood sugar'}>],
  [<Axes: title={'center': 'resting ecg'}>,
   <Axes: title={'center': 'max heart rate'}>,
   <Axes: title={'center': 'exercise angina'}>],
  [<Axes: title={'center': 'oldpeak'}>,
   <Axes: title={'center': 'ST slope'}>,
   <Axes: title={'center': 'target'}>]], dtype=object)
```



```
In [13]: (df == 0.0).sum()
```

```
Out[13]: age          0
          sex         281
          chest pain type  0
          resting bp s    1
          cholesterol     172
          fasting blood sugar 936
          resting ecg      684
          max heart rate   0
          exercise angina  729
          oldpeak          455
          ST slope          1
          target            561
          dtype: int64
```

## Handling Missing/Zero Values

Replaces 0 values in cholesterol with the median of non-zero cholesterol values.

This avoids biasing the model with invalid zeros.

```
In [14]: median_cholesterol = df[df['cholesterol'] > 0]['cholesterol'].median()
print(median_cholesterol)
```

240.0

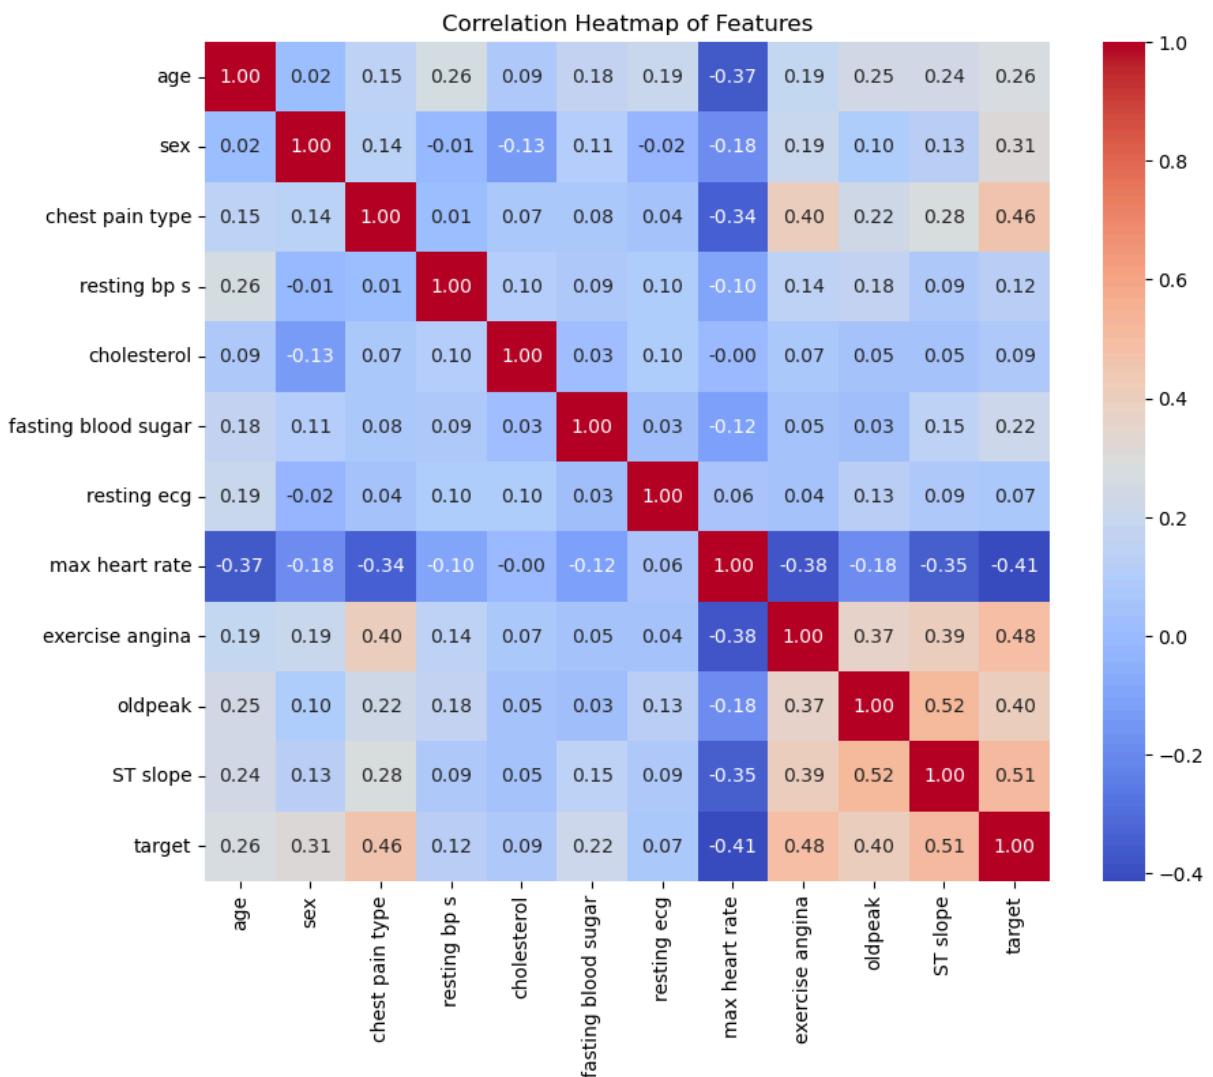
```
In [15]: df['cholesterol'] = df['cholesterol'].replace(0, median_cholesterol)
```

finding data where their values is 0.0

```
In [16]: (df == 0.0).sum()
```

```
Out[16]: age          0  
sex         281  
chest pain type    0  
resting bp s        1  
cholesterol        0  
fasting blood sugar 936  
resting ecg         684  
max heart rate      0  
exercise angina     729  
oldpeak            455  
ST slope           1  
target             561  
dtype: int64
```

```
In [17]: plt.figure(figsize=(10,8))  
sns.heatmap(df.corr(), annot=True, cmap="coolwarm", fmt=".2f")  
plt.title("Correlation Heatmap of Features")  
plt.show()
```



## Dropping Less Useful Features ie, negative correlated columns

Keeps the dataset clean and reduces noise.

```
In [18]: df1 = df.drop(['resting bp s', 'cholesterol',
                     'fasting blood sugar', 'resting ecg', 'max heart rate'], axis=1)
```

```
In [19]: df1
```

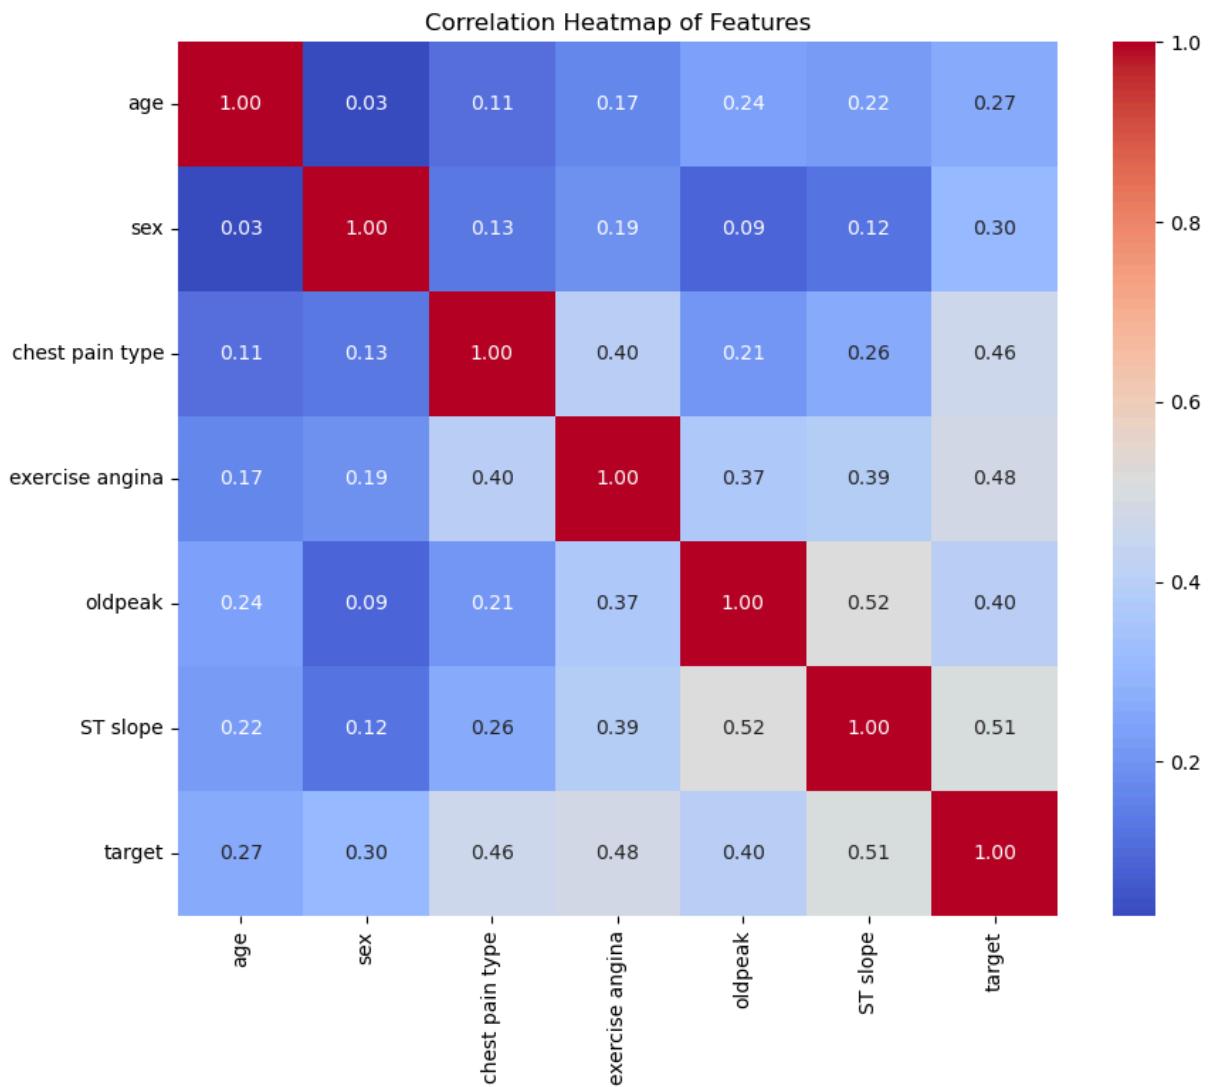
Out[19]:

	age	sex	chest pain type	exercise angina	oldpeak	ST slope	target
0	40	1	2	0	0.0	1	0
1	49	0	3	0	1.0	2	1
2	37	1	2	0	0.0	1	0
3	48	0	4	1	1.5	2	1
4	54	1	3	0	0.0	1	0
...	...	...	...	...	...	...	...
1185	45	1	1	0	1.2	2	1
1186	68	1	4	0	3.4	2	1
1187	57	1	4	1	1.2	2	1
1188	57	0	2	0	0.0	2	1
1189	38	1	3	0	0.0	1	0

1190 rows × 7 columns

In [20]: `df2 = df1[(df1['age'] >= 35) & (df1['age'] <= 70)]`

In [21]: `plt.figure(figsize=(10,8))  
sns.heatmap(df2.corr(), annot=True, cmap="coolwarm", fmt=".2f")  
plt.title("Correlation Heatmap of Features")  
plt.show()`



## Checking Data Distribution

Prints the distribution of values for each column.

Helps understand categorical feature balance.

```
In [22]: for col in df2.columns:
    print(f"\nColumn: {col}")
    print(df2[col].value_counts())
```

Column: age

age

54	67
58	58
57	50
52	47
55	47
59	47
56	47
51	47
62	46
60	44
53	40
61	38
48	38
63	37
41	33
50	32
43	31
64	31
46	31
65	29
44	29
49	27
42	26
45	25
47	23
67	23
66	19
39	18
38	17
69	16
40	16
35	14
68	13
37	13
70	11
36	6

Name: count, dtype: int64

Column: sex

sex

1	873
0	263

Name: count, dtype: int64

Column: chest pain type

chest pain type

4	608
3	271
2	197
1	60

Name: count, dtype: int64

Column: exercise angina

exercise angina

```
0      691
1      445
Name: count, dtype: int64

Column: oldpeak
oldpeak
 0.0    429
 1.0     96
 2.0     82
 1.5     55
 1.2     40
 3.0     31
 1.4     30
 0.2     29
 0.8     27
 0.6     26
 1.8     25
 1.6     24
 0.5     24
 0.1     20
 2.5     18
 0.4     17
 0.3     14
 2.6     13
 1.9     12
 2.8     11
 2.2      9
 4.0      9
 3.6      8
 2.4      7
 1.3      7
 0.9      7
 1.1      7
 1.7      6
 3.4      5
 0.7      5
 3.2      4
 4.2      4
 2.3      4
 2.1      3
 5.6      2
 2.9      2
 3.5      2
 6.2      2
 3.1      2
 3.8      2
 -0.1     2
 -0.5     2
 -1.0     2
 -2.0     1
  3.7     1
 -0.8     1
 -0.7     1
 -1.1     1
 -1.5     1
 -2.6     1
```

```
-0.9      1
5.0      1
4.4      1
Name: count, dtype: int64

Column: ST slope
ST slope
2      564
1      494
3      77
0      1
Name: count, dtype: int64

Column: target
target
1      606
0      530
Name: count, dtype: int64
```

In [23]: `print(df2.dtypes)`

```
age          int64
sex          int64
chest pain type  int64
exercise angina  int64
oldpeak      float64
ST slope      int64
target        int64
dtype: object
```

## Train-Test Split

Splits the data into training set (80%) and testing set (20%)

In [24]: `from sklearn.model_selection import train_test_split`

```
X = df2.drop('target', axis=1)
y = df2['target']

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=45, stratify=y # stratify keeps balance
)

print("Train target distribution:\n", y_train.value_counts())
print("Test target distribution:\n", y_test.value_counts())
```

```
Train target distribution:  
target  
1    484  
0    424  
Name: count, dtype: int64  
Test target distribution:  
target  
1    122  
0    106  
Name: count, dtype: int64
```

```
In [25]: for col in ['chest pain type', 'exercise angina', 'ST slope','oldpeak']:  
    print(f"\n{col} - Train:")  
    print(X_train[col].value_counts(normalize=True))  
    print(f"{col} - Test:")  
    print(X_test[col].value_counts(normalize=True))
```

```
chest pain type - Train:  
chest pain type  
4    0.530837  
3    0.245595  
2    0.168502  
1    0.055066  
Name: proportion, dtype: float64  
chest pain type - Test:  
chest pain type  
4    0.552632  
3    0.210526  
2    0.192982  
1    0.043860  
Name: proportion, dtype: float64  
  
exercise angina - Train:  
exercise angina  
0    0.602423  
1    0.397577  
Name: proportion, dtype: float64  
exercise angina - Test:  
exercise angina  
0    0.631579  
1    0.368421  
Name: proportion, dtype: float64  
  
ST slope - Train:  
ST slope  
2    0.498899  
1    0.435022  
3    0.066079  
Name: proportion, dtype: float64  
ST slope - Test:  
ST slope  
2    0.486842  
1    0.434211  
3    0.074561  
0    0.004386  
Name: proportion, dtype: float64  
  
oldpeak - Train:  
oldpeak  
0.0    0.386564  
1.0    0.087004  
2.0    0.073789  
1.5    0.046256  
1.2    0.030837  
1.4    0.028634  
3.0    0.026432  
0.2    0.026432  
1.8    0.023128  
0.8    0.023128  
1.6    0.020925  
0.5    0.019824  
0.6    0.019824  
0.1    0.018722
```

```
0.4    0.016520
2.5    0.016520
0.3    0.012115
2.6    0.012115
1.9    0.009912
2.8    0.009912
2.2    0.008811
1.7    0.006608
1.1    0.006608
3.6    0.006608
1.3    0.006608
4.0    0.005507
0.7    0.005507
0.9    0.005507
2.4    0.004405
4.2    0.003304
3.1    0.002203
2.9    0.002203
-0.1   0.002203
5.6    0.002203
3.5    0.002203
6.2    0.002203
3.8    0.002203
2.3    0.002203
3.2    0.002203
-1.0   0.002203
-0.9   0.001101
-0.5   0.001101
3.7    0.001101
5.0    0.001101
-2.0   0.001101
-1.5   0.001101
4.4    0.001101
2.1    0.001101
-0.7   0.001101
```

Name: proportion, dtype: float64

oldpeak - Test:

```
oldpeak
0.0    0.342105
1.0    0.074561
2.0    0.065789
1.5    0.057018
1.2    0.052632
0.6    0.035088
3.0    0.030702
0.5    0.026316
0.8    0.026316
3.4    0.021930
0.2    0.021930
1.6    0.021930
1.4    0.017544
1.8    0.017544
4.0    0.017544
1.9    0.013158
0.3    0.013158
0.1    0.013158
```

```

2.4    0.013158
2.5    0.013158
3.6    0.008772
0.4    0.008772
3.2    0.008772
0.9    0.008772
2.6    0.008772
2.3    0.008772
2.1    0.008772
2.8    0.008772
2.2    0.004386
-0.5   0.004386
4.2    0.004386
-1.1   0.004386
1.1    0.004386
1.3    0.004386
-2.6   0.004386
-0.8   0.004386
Name: proportion, dtype: float64

```

## Handling Negative Values

Replaces negative values in oldpeak with 0.

Negative values are invalid in this context.

```
In [26]: X['oldpeak'] = X['oldpeak'].apply(lambda x: 0 if x < 0 else x)
X_train['oldpeak'] = X_train['oldpeak'].apply(lambda x: 0 if x < 0 else x)
X_test['oldpeak'] = X_test['oldpeak'].apply(lambda x: 0 if x < 0 else x)
```

```
In [27]: print(df2['oldpeak'])
```

```

0      0.0
1      1.0
2      0.0
3      1.5
4      0.0
...
1185   1.2
1186   3.4
1187   1.2
1188   0.0
1189   0.0
Name: oldpeak, Length: 1136, dtype: float64

```

```
In [28]: for col in ['oldpeak']:
    print(f"\n{col} - Train:")
    print(X_train[col].value_counts(normalize=True))
    print(f"{col} - Test:")
    print(X_test[col].value_counts(normalize=True))
```

oldpeak - Train:

oldpeak  
0.0 0.396476  
1.0 0.087004  
2.0 0.073789  
1.5 0.046256  
1.2 0.030837  
1.4 0.028634  
3.0 0.026432  
0.2 0.026432  
0.8 0.023128  
1.8 0.023128  
1.6 0.020925  
0.6 0.019824  
0.5 0.019824  
0.1 0.018722  
0.4 0.016520  
2.5 0.016520  
2.6 0.012115  
0.3 0.012115  
1.9 0.009912  
2.8 0.009912  
2.2 0.008811  
1.7 0.006608  
1.3 0.006608  
3.6 0.006608  
1.1 0.006608  
4.0 0.005507  
0.7 0.005507  
0.9 0.005507  
2.4 0.004405  
4.2 0.003304  
6.2 0.002203  
3.1 0.002203  
2.3 0.002203  
3.8 0.002203  
3.5 0.002203  
5.6 0.002203  
3.2 0.002203  
2.9 0.002203  
4.4 0.001101  
3.7 0.001101  
2.1 0.001101  
5.0 0.001101

Name: proportion, dtype: float64

oldpeak - Test:

oldpeak  
0.0 0.359649  
1.0 0.074561  
2.0 0.065789  
1.5 0.057018  
1.2 0.052632  
0.6 0.035088  
3.0 0.030702  
0.5 0.026316  
0.8 0.026316

```

3.4    0.021930
0.2    0.021930
1.6    0.021930
1.4    0.017544
1.8    0.017544
4.0    0.017544
1.9    0.013158
0.1    0.013158
0.3    0.013158
2.4    0.013158
2.5    0.013158
3.6    0.008772
0.4    0.008772
3.2    0.008772
2.6    0.008772
0.9    0.008772
2.3    0.008772
2.1    0.008772
2.8    0.008772
2.2    0.004386
4.2    0.004386
1.1    0.004386
1.3    0.004386
Name: proportion, dtype: float64

```

```
In [29]: df2_train = X_train.copy()
df2_train['target'] = y_train
df2_test = X_test.copy()
df2_test['target'] = y_test
```

## Comparing Feature Distribution in Train vs Test

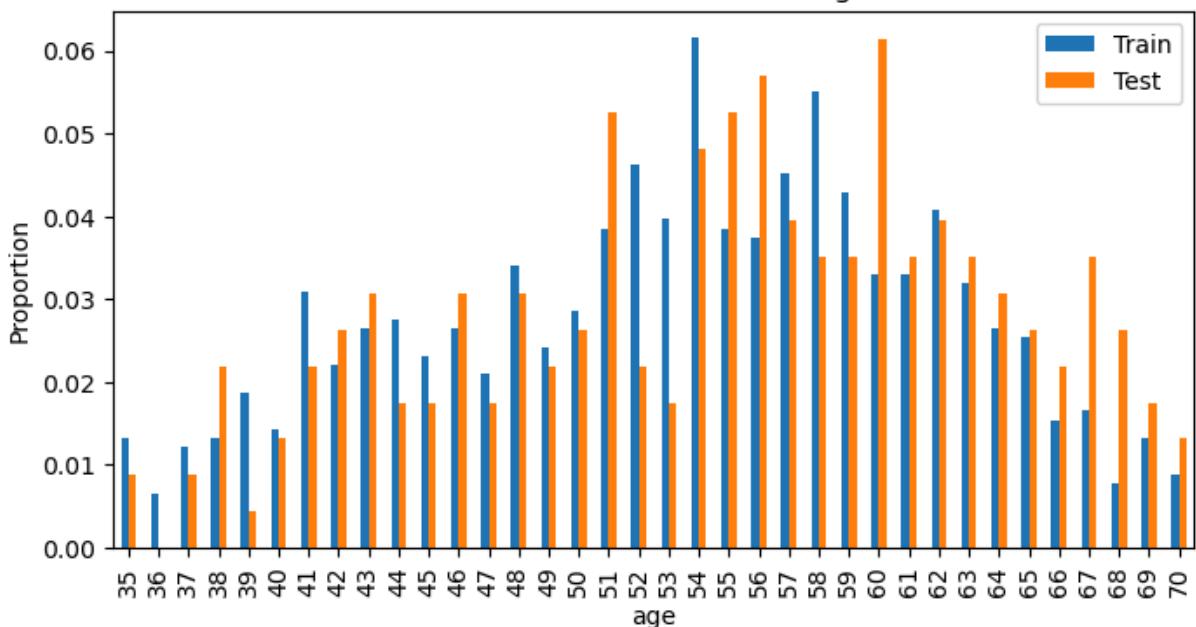
Plots bar charts for selected categorical features (sex, chest pain type, exercise angina, ST slope) to verify that train and test distributions are similar.

```
In [30]: cols = ['age', 'sex', 'chest pain type', 'exercise angina', 'ST slope']

for col in cols:
    plt.figure(figsize=(8,4))
    train_counts = df2_train[col].value_counts(normalize=True).sort_index()
    test_counts = df2_test[col].value_counts(normalize=True).sort_index()
    df_plot = pd.DataFrame({'Train': train_counts, 'Test': test_counts})
    df_plot.plot(kind='bar', figsize=(8,4))
    plt.title(f'Train vs Test Distribution: {col}')
    plt.ylabel('Proportion')
    plt.xlabel(col)
    plt.show()
```

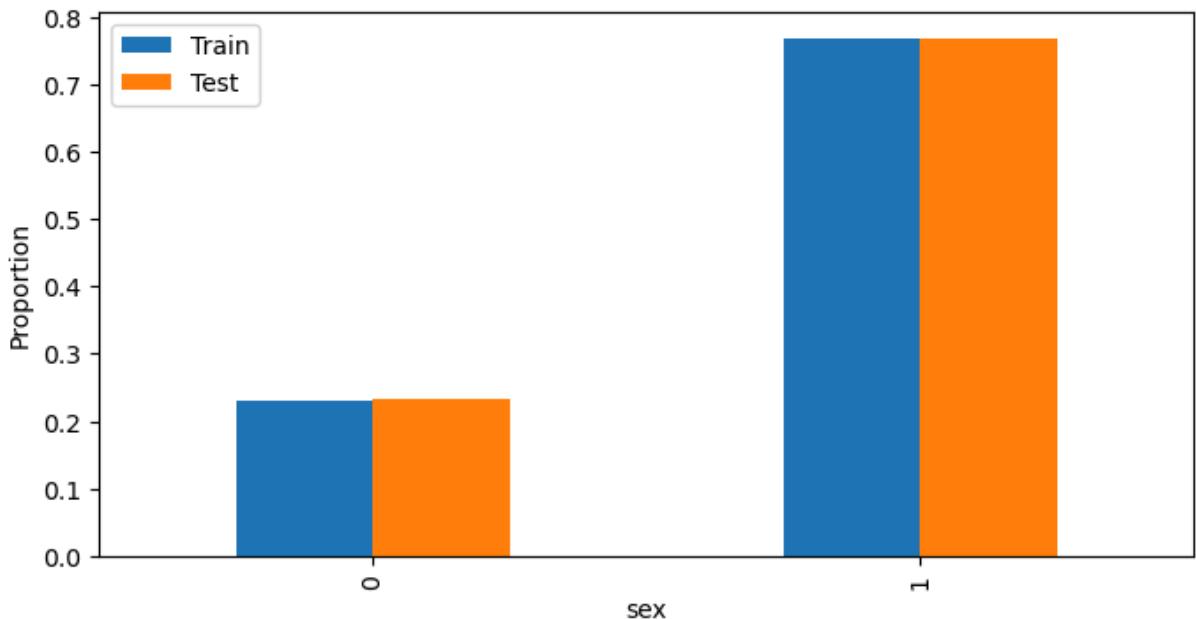
<Figure size 800x400 with 0 Axes>

## Train vs Test Distribution: age



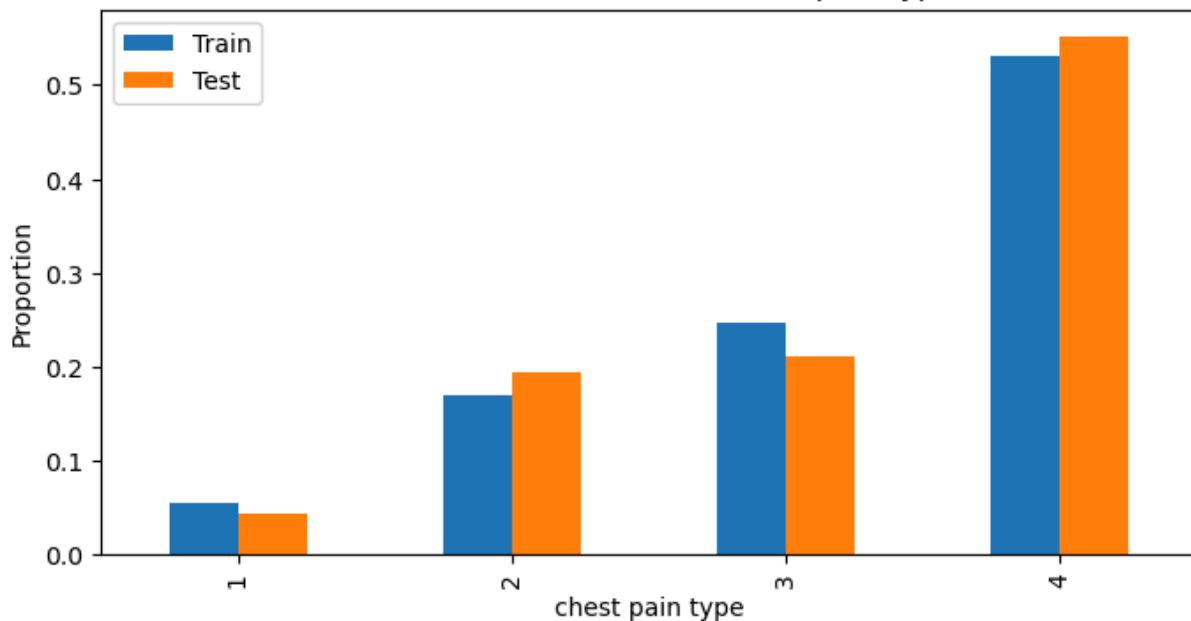
&lt;Figure size 800x400 with 0 Axes&gt;

## Train vs Test Distribution: sex



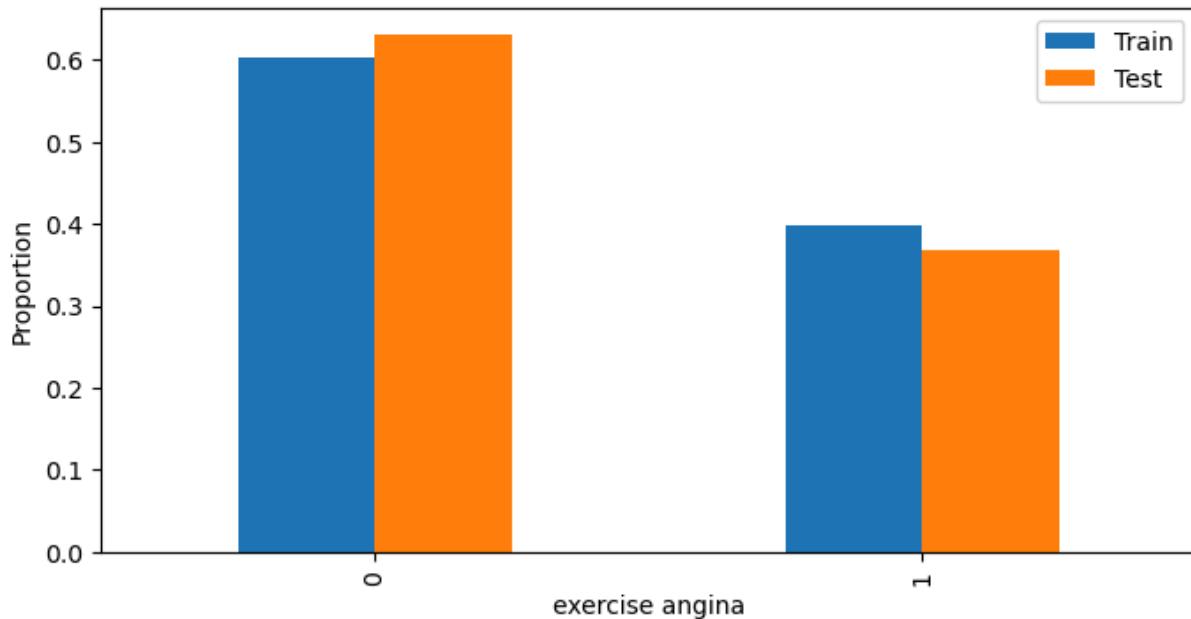
&lt;Figure size 800x400 with 0 Axes&gt;

Train vs Test Distribution: chest pain type



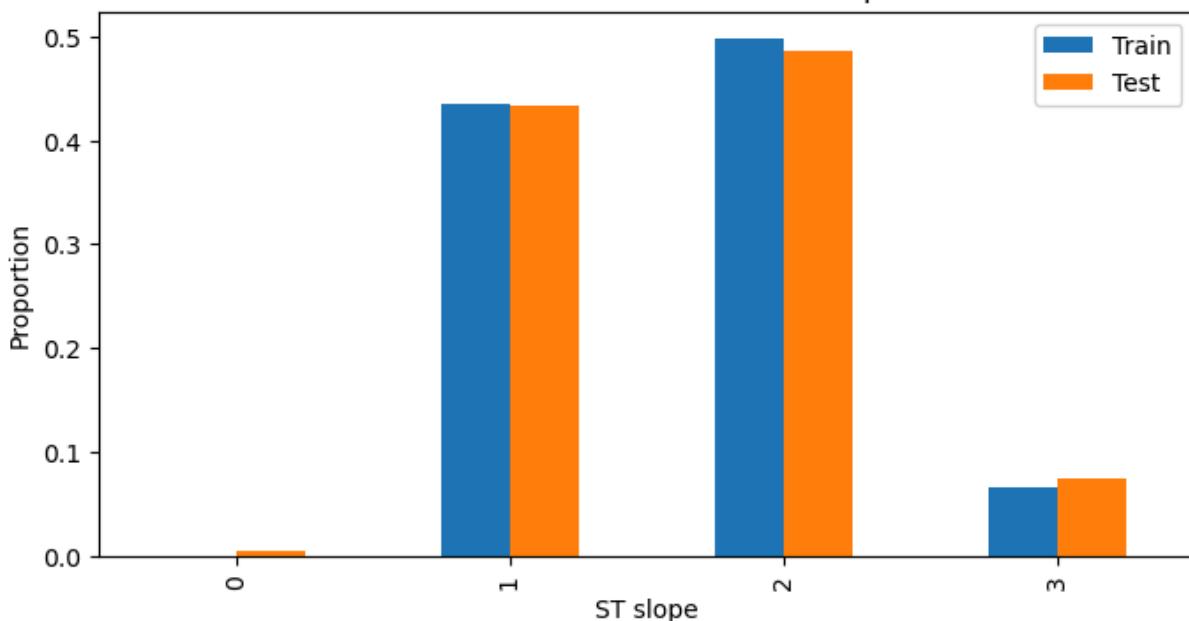
&lt;Figure size 800x400 with 0 Axes&gt;

Train vs Test Distribution: exercise angina



&lt;Figure size 800x400 with 0 Axes&gt;

Train vs Test Distribution: ST slope



## Train Logistic Regression Model

Create a pipeline with feature scaling and logistic regression, then evaluate performance.

```
In [31]: from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix

pipeline = Pipeline([
    ('scaler', StandardScaler()),
    ('model', LogisticRegression())
])

# Train
pipeline.fit(X_train, y_train)

# Predict
y_pred = pipeline.predict(X_test)

# Evaluate
print("Accuracy:", accuracy_score(y_test, y_pred))
```

Accuracy: 0.8421052631578947

## Decision Tree Model

Train and evaluate a decision tree classifier on selected features.

```
In [32]: from sklearn.tree import DecisionTreeClassifier

# Pipeline
```

```

pipe_dt = Pipeline([
    ('scaler', StandardScaler()),
    ('dt', DecisionTreeClassifier(random_state=45))
])

# Fit
pipe_dt.fit(X_train[cols], y_train)

# Predict
y_pred_dt = pipe_dt.predict(X_test[cols])

# Evaluate
print("Decision Tree Accuracy:", accuracy_score(y_test, y_pred_dt))

```

Decision Tree Accuracy: 0.8289473684210527

## Model 3: Random Forest Classifier

Ensemble approach using 1000 decision trees to improve prediction accuracy and robustness.

```

In [33]: from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report

# Random Forest pipeline
pipe_rf = Pipeline([
    ('scaler', StandardScaler()), # optional for tree-based models
    ('rf', RandomForestClassifier(n_estimators=1000, random_state=45))
])

# Train
pipe_rf.fit(X_train, y_train)

# Predict
y_pred_rf = pipe_rf.predict(X_test)

# Evaluate
print("Random Forest Accuracy:", accuracy_score(y_test, y_pred_rf))

print("\nClassification Report:\n", classification_report(y_test, y_pred_rf))

```

Random Forest Accuracy: 0.9035087719298246

Classification Report:				
	precision	recall	f1-score	support
0	0.90	0.89	0.90	106
1	0.90	0.92	0.91	122
accuracy			0.90	228
macro avg	0.90	0.90	0.90	228
weighted avg	0.90	0.90	0.90	228

## Model 4: Linear Regression (Regression → Classification)

Converting a regression problem to classification by thresholding continuous predictions at 0.5.

```
In [34]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

# Linear Regression pipeline
pipe_lr_model = Pipeline([
    ('scaler', StandardScaler()),
    ('lr', LinearRegression())
])

# Train
pipe_lr_model.fit(X_train, y_train)

# Predict
y_pred_lr = pipe_lr_model.predict(X_test)

# Evaluate

print("MSE:", mean_squared_error(y_test, y_pred_lr))
print("R2 Score:", r2_score(y_test, y_pred_lr))

y_pred_lr_bin = (y_pred_lr >= 0.5).astype(int)
from sklearn.metrics import accuracy_score
print("Linear Regression Accuracy (binary):", accuracy_score(y_test, y_pred_lr))

MSE: 0.1344726329910612
R2 Score: 0.4594474672589449
Linear Regression Accuracy (binary): 0.8333333333333334
```

## K-Nearest Neighbors Model

Distance-based classifier using the 5 nearest neighbors for prediction.

```
In [35]: from sklearn.neighbors import KNeighborsClassifier

# KNN pipeline
pipe_knn = Pipeline([
    ('scaler', StandardScaler()), # scaling is important for KNN
    ('knn', KNeighborsClassifier(n_neighbors=5)) # you can tune n_neighbors
])

# Train
pipe_knn.fit(X_train, y_train)

# Predict
y_pred_knn = pipe_knn.predict(X_test)
```

```
# Evaluate
print("KNN Accuracy:", accuracy_score(y_test, y_pred_knn))
```

KNN Accuracy: 0.8464912280701754

## Gradient Boosting Model

Boosting ensemble with 1000 trees, learning rate 0.10, and detailed performance evaluation.

```
In [36]: from sklearn.ensemble import GradientBoostingClassifier

# Initialize the Gradient Boosting model
gb_model = GradientBoostingClassifier(
    n_estimators=1000,      # number of trees
    learning_rate=0.10,     # step size shr
    max_depth=4,           # depth of each tree
    random_state=45
)

# Train the model
gb_model.fit(X_train, y_train)

# Make predictions
y_pred_gb = gb_model.predict(X_test)

# Evaluate performance
print("Gradient Boosting Accuracy:", accuracy_score(y_test, y_pred_gb))
print("\nClassification Report:\n", classification_report(y_test, y_pred_gb))
```

Gradient Boosting Accuracy: 0.9035087719298246

Classification Report:

	precision	recall	f1-score	support
0	0.91	0.88	0.89	106
1	0.90	0.93	0.91	122
accuracy			0.90	228
macro avg	0.90	0.90	0.90	228
weighted avg	0.90	0.90	0.90	228

## Model Comparison & Evaluation

Systematically comparing all models using multiple metrics:

- Accuracy: Overall correctness
- Precision: Positive prediction accuracy
- Recall: True positive detection rate
- F1-Score: Harmonic mean of precision and recall

```
In [37]: import pandas as pd
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

# Create a dictionary to store your models and predictions
models_preds = {
    'Logistic Regression': y_pred,
    'Decision Tree': y_pred_dt,
    'Random Forest': y_pred_rf,
    'Linear Regression (binary)': y_pred_lr_bin,
    'K-Nearest Neighbors': y_pred_knn,
    'Gradient Boosting': y_pred_gb
}

# Initialize a list to store metrics
metrics_list = []

# Calculate metrics for each model
for model_name, preds in models_preds.items():
    accuracy = accuracy_score(y_test, preds)
    precision = precision_score(y_test, preds)
    recall = recall_score(y_test, preds)
    f1 = f1_score(y_test, preds)

    metrics_list.append({
        'Model': model_name,
        'Accuracy': round(accuracy, 3),
        'Precision': round(precision, 3),
        'Recall': round(recall, 3),
        'F1-score': round(f1, 3),
    })

# Convert to DataFrame
metrics_df = pd.DataFrame(metrics_list)
metrics_df = metrics_df.sort_values(by='Accuracy', ascending=False).reset_index()

print(metrics_df)
```

	Model	Accuracy	Precision	Recall	F1-score
0	Random Forest	0.904	0.903	0.918	0.911
1	Gradient Boosting	0.904	0.897	0.926	0.911
2	K-Nearest Neighbors	0.846	0.848	0.869	0.858
3	Logistic Regression	0.842	0.831	0.885	0.857
4	Linear Regression (binary)	0.833	0.828	0.869	0.848
5	Decision Tree	0.829	0.867	0.803	0.834

## Model Performance Visualizations

- 1. Accuracy Bar Chart:** Side-by-side comparison of all models
- 2. ROC Curves:** Performance analysis for models with probability predictions (AUC scores)

```
In [38]: import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import roc_curve, auc, confusion_matrix, classification_report

# --- 1. Accuracy Table and Bar Plot ---
accuracy_dict = {
    'Logistic Regression': accuracy_score(y_test, y_pred),
    'Decision Tree': accuracy_score(y_test, y_pred_dt),
    'Random Forest': accuracy_score(y_test, y_pred_rf),
    'Linear Regression (binary)': accuracy_score(y_test, y_pred_lr_bin),
    'K-Nearest Neighbors': accuracy_score(y_test, y_pred_knn),
    'Gradient Boosting': accuracy_score(y_test, y_pred_gb)
}

accuracy_table = pd.DataFrame(list(accuracy_dict.items()), columns=['Model'],
print(accuracy_table)

# Bar plot
plt.figure(figsize=(8,5))
sns.barplot(x='Accuracy', y='Model', data=accuracy_table, palette='viridis')
plt.title('Model Accuracy Comparison')
plt.xlim(0,1)
plt.show()

# --- 3. ROC Curves ---
models_proba = {
    'Logistic Regression': pipeline,
    'Random Forest': pipe_rf,
    'Gradient Boosting': gb_model,
    'K-Nearest Neighbors': pipe_knn
}

plt.figure(figsize=(8,6))
for name, model in models_proba.items():
    try:
        y_proba = model.predict_proba(X_test)[:,1]
        fpr, tpr, _ = roc_curve(y_test, y_proba)
        roc_auc = auc(fpr, tpr)
        plt.plot(fpr, tpr, label=f'{name} (AUC={roc_auc:.2f})')
    except:
        print(f'{name} skipped (no predict_proba)')

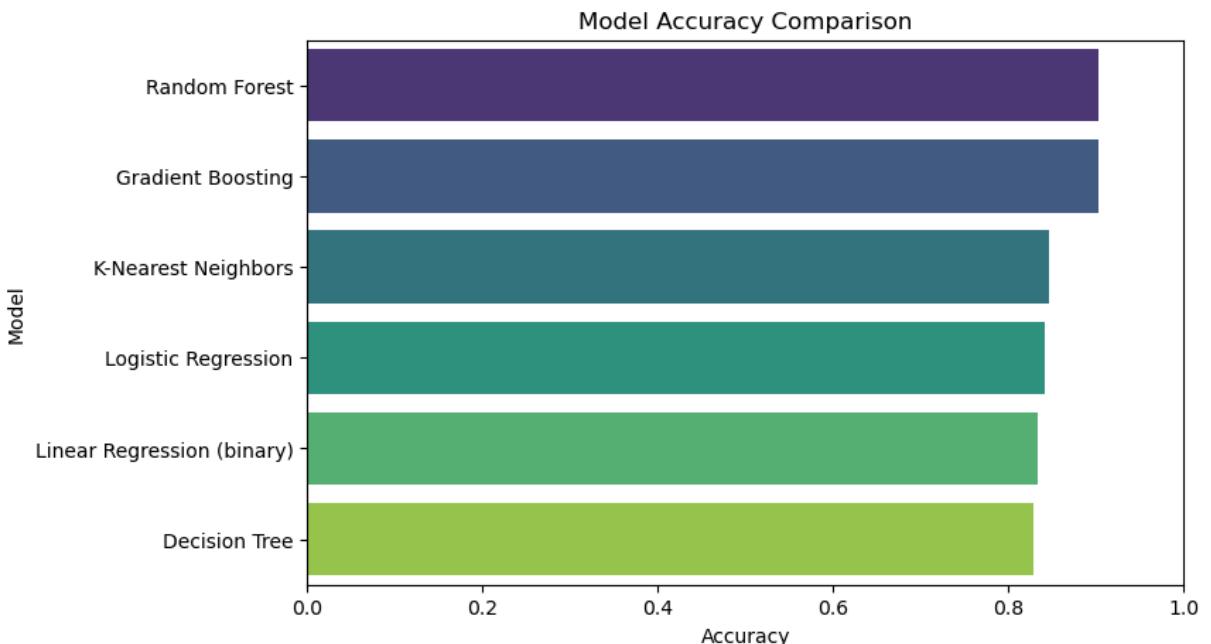
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curves')
plt.legend()
plt.grid(True)
plt.show()
```

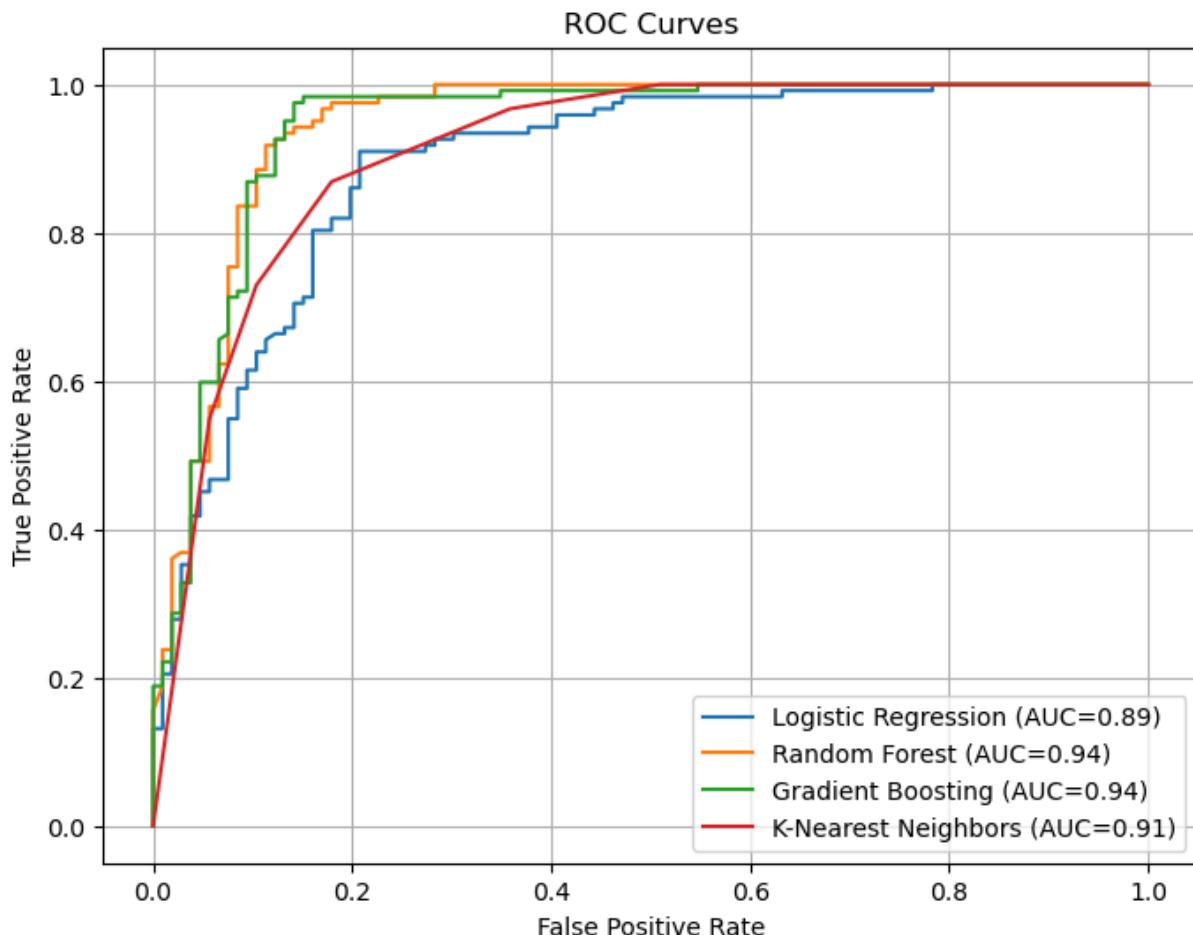
	Model	Accuracy
2	Random Forest	0.903509
5	Gradient Boosting	0.903509
4	K-Nearest Neighbors	0.846491
0	Logistic Regression	0.842105
3	Linear Regression (binary)	0.833333
1	Decision Tree	0.828947

/var/folders/d6/jbdgpdt2dlc2370hmzb47k0000gn/T/ipykernel\_16597/894066980.py:20: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x='Accuracy', y='Model', data=accuracy_table, palette='viridis')
```





## Confusion Matrix Analysis

Detailed breakdown of model predictions vs actual results:

- **True Positives (TP):** Correctly predicted disease cases
- **True Negatives (TN):** Correctly predicted no disease cases
- **False Positives (FP):** Incorrectly predicted disease (Type I error)
- **False Negatives (FN):** Missed disease cases (Type II error)

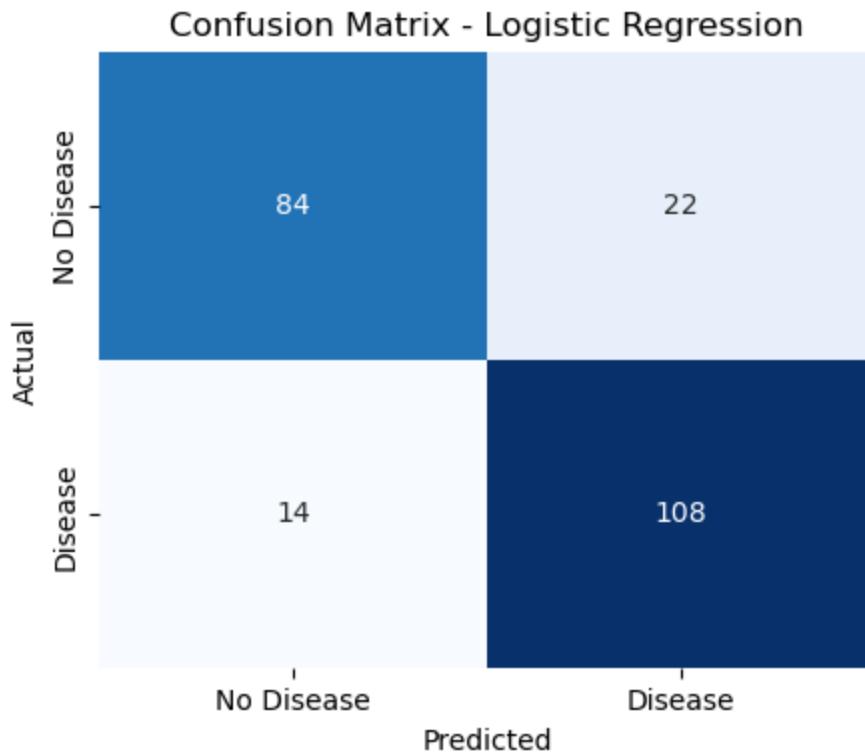
Visualizing confusion matrices for all 6 models to understand prediction patterns and error types.

```
In [39]: from sklearn.metrics import confusion_matrix

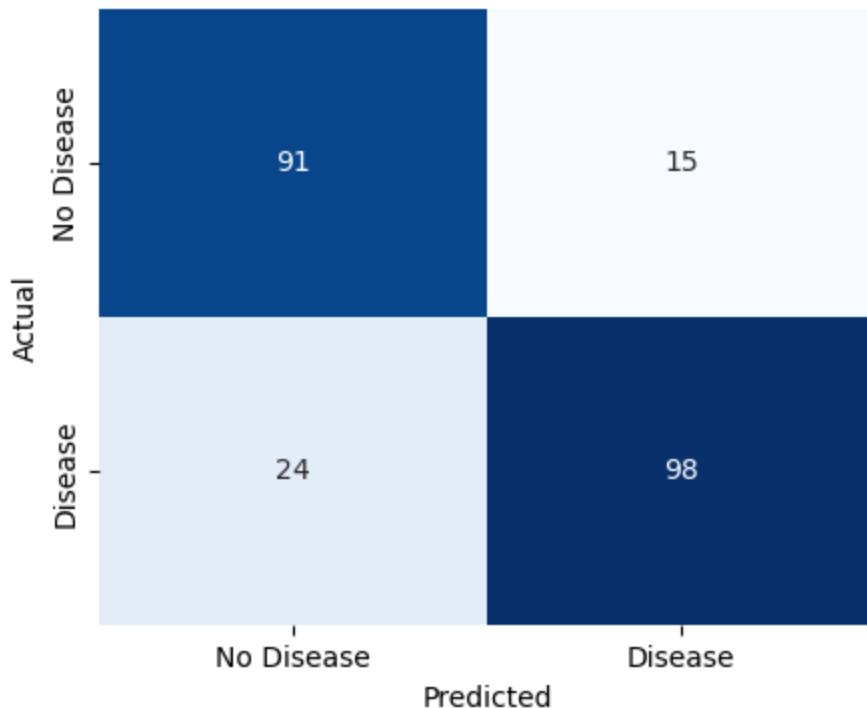
# Function to plot confusion matrix
def plot_conf_matrix(y_true, y_pred, model_name):
    cm = confusion_matrix(y_true, y_pred)
    plt.figure(figsize=(5,4))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=False,
                xticklabels=['No Disease', 'Disease'],
                yticklabels=['No Disease', 'Disease'])
    plt.title(f'Confusion Matrix - {model_name}')
    plt.xlabel('Predicted')
```

```
plt.ylabel('Actual')
plt.show()

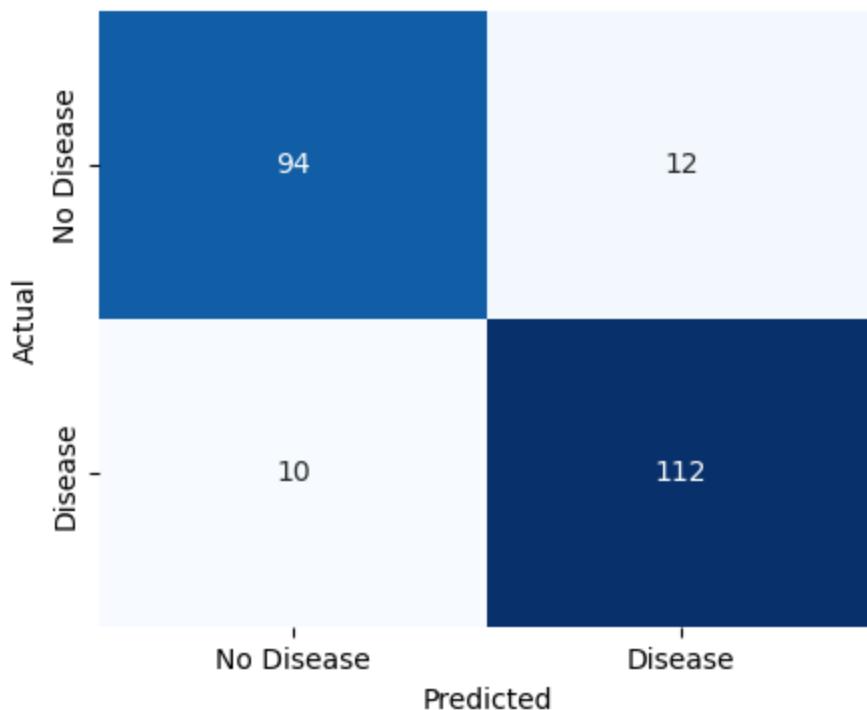
# Example usage for your models:
plot_conf_matrix(y_test, y_pred, "Logistic Regression")
plot_conf_matrix(y_test, y_pred_dt, "Decision Tree")
plot_conf_matrix(y_test, y_pred_rf, "Random Forest")
plot_conf_matrix(y_test, y_pred_lr_bin, "Linear Regression (Binary)")
plot_conf_matrix(y_test, y_pred_knn, "KNN")
plot_conf_matrix(y_test, y_pred_gb, "Gradient Boosting")
```

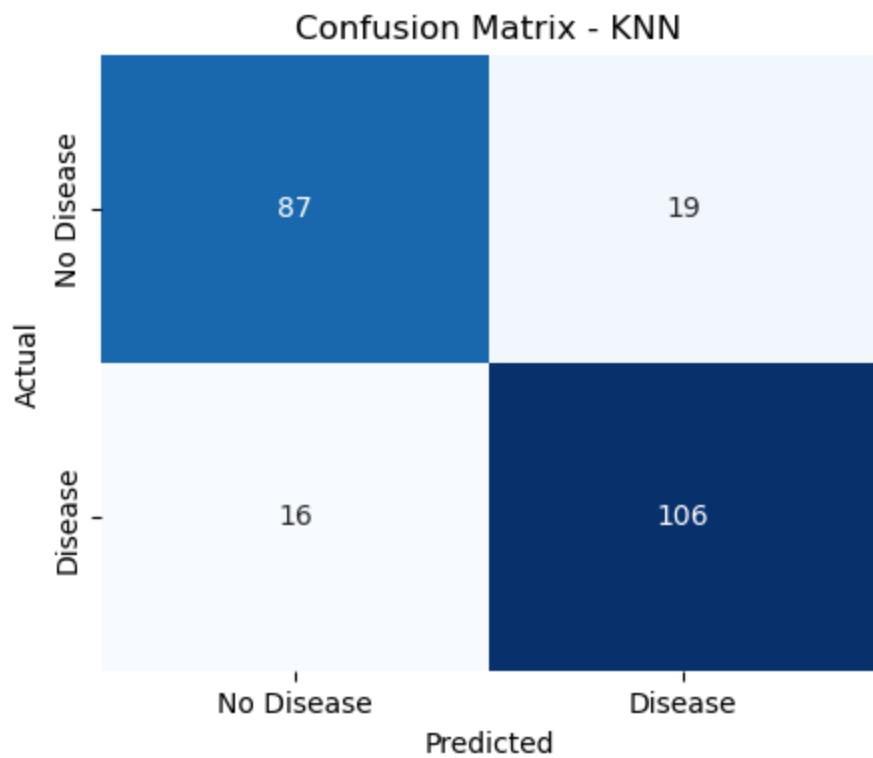
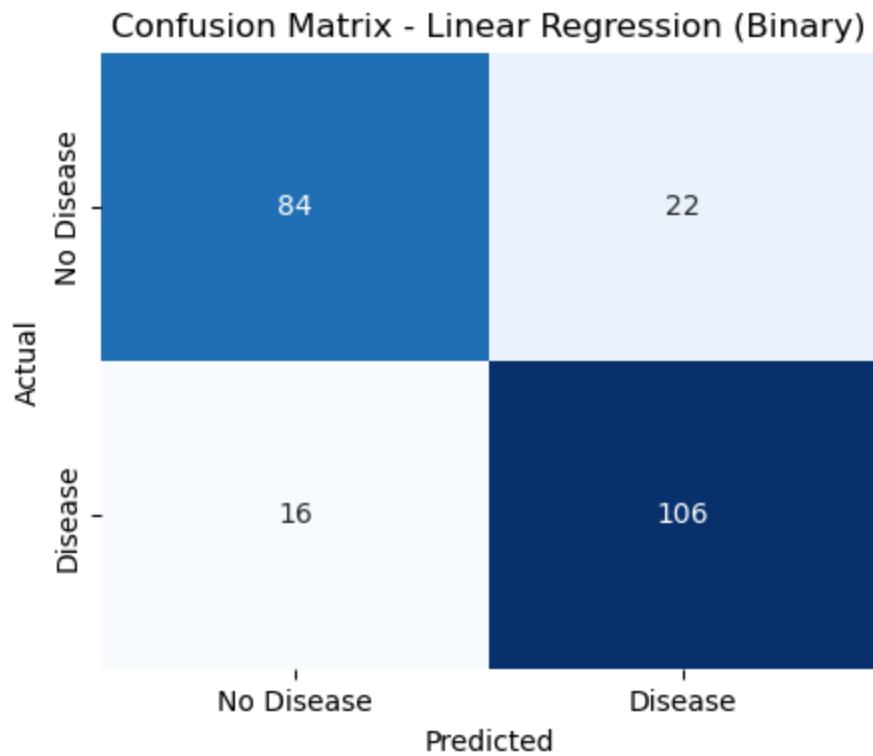


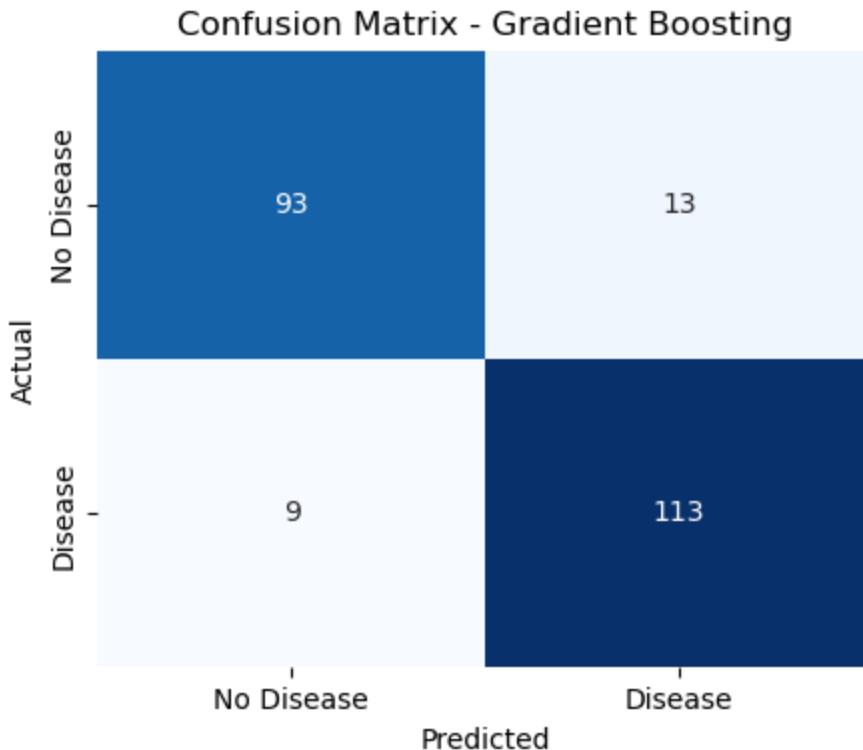
Confusion Matrix - Decision Tree



Confusion Matrix - Random Forest







## Scatter Plot: Age vs Chest Pain Type

Visualizing how age and chest pain type relate to heart disease diagnosis.

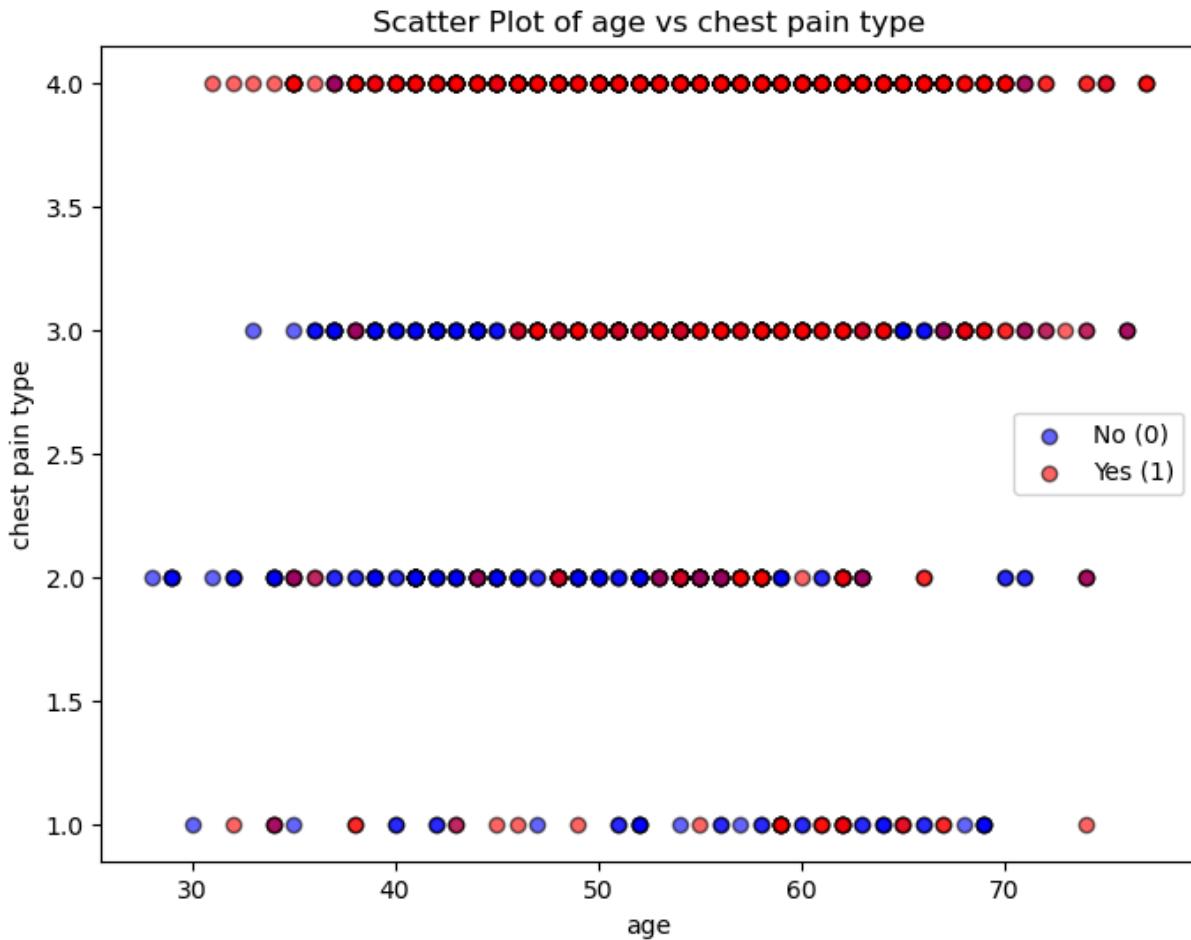
```
In [40]: import matplotlib.pyplot as plt

x_feature = 'age'
y_feature = 'chest pain type'

# Scatter plot
plt.figure(figsize=(8,6))
plt.scatter(df[df['target']==0][x_feature], df[df['target']==0][y_feature],
            color='blue', label='No (0)', alpha=0.6, edgecolor='k')

plt.scatter(df[df['target']==1][x_feature], df[df['target']==1][y_feature],
            color='red', label='Yes (1)', alpha=0.6, edgecolor='k')

plt.xlabel(x_feature)
plt.ylabel(y_feature)
plt.title(f"Scatter Plot of {x_feature} vs {y_feature}")
plt.legend()
plt.show()
```



```
In [41]: X_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 908 entries, 493 to 150
Data columns (total 6 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   age              908 non-null    int64  
 1   sex              908 non-null    int64  
 2   chest pain type  908 non-null    int64  
 3   exercise angina  908 non-null    int64  
 4   oldpeak          908 non-null    float64 
 5   ST slope         908 non-null    int64  
dtypes: float64(1), int64(5)
memory usage: 49.7 KB
```

## Model Export & Feature Importance

- Save Model:** Export trained Gradient Boosting classifier for future use
- Save Baseline:** Store healthy patient averages as reference
- Feature Ranking:** Display importance scores for each predictor

```
In [42]: # Import libraries
import pandas as pd
import numpy as np
```

```
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.model_selection import train_test_split
import joblib
import json
# Save the newly trained model
joblib.dump(gb_model, "model.pkl")
print("Model saved!")

# Save healthy averages
features_needed = ['age', 'sex', 'chest pain type', 'exercise angina', 'oldpeak']
healthy_avg = df[df['target'] == 0][features_needed].mean()

with open('healthy_avg.json', 'w') as f:
    json.dump([healthy_avg.to_dict()], f, indent=4)
print("Healthy averages saved!")

# Check if model has feature_importances_
loaded_model = joblib.load("model.pkl")
if hasattr(loaded_model, 'feature_importances_'):
    print("\nFeature importances available")
    feature_names = ['Age', 'Sex', 'Chest Pain Type', 'Exercise Angina', 'Old Peak', 'ST Slope']
    for name, importance in zip(feature_names, loaded_model.feature_importances_):
        print(f"  {name}: {importance:.4f}")
else:
    print("\n Model doesn't have feature_importances_")
```

Model saved!

Healthy averages saved!

Feature importances available  
Age: 0.1593  
Sex: 0.0617  
Chest Pain Type: 0.1764  
Exercise Angina: 0.0522  
Oldpeak: 0.1562  
ST Slope: 0.3941

In [ ]: